



UNIVERSITAT DE  
BARCELONA

# Master in Fundamental Principles of Data Science

Dr Rohit Kumar



UNIVERSITAT DE  
BARCELONA

# Hadoop

# What is Hadoop


Apache Hadoop is an open source software framework used to develop data processing applications which are executed in a distributed computing environment.

In Hadoop, data resides in a distributed file system which is called as a **Hadoop Distributed File system**.

The processing model is based on '**Data Locality**' concept wherein computational logic is sent to cluster nodes(server) containing data.



# Hadoop



**Distributed  
Computation**

A new programming model  
called **MapReduce**  
Is used to run computation.

**Distributed  
Storage**

**HDFS** a distributed file  
system provides  
storage

Server cluster



Hadoop can run on commodity  
hardware

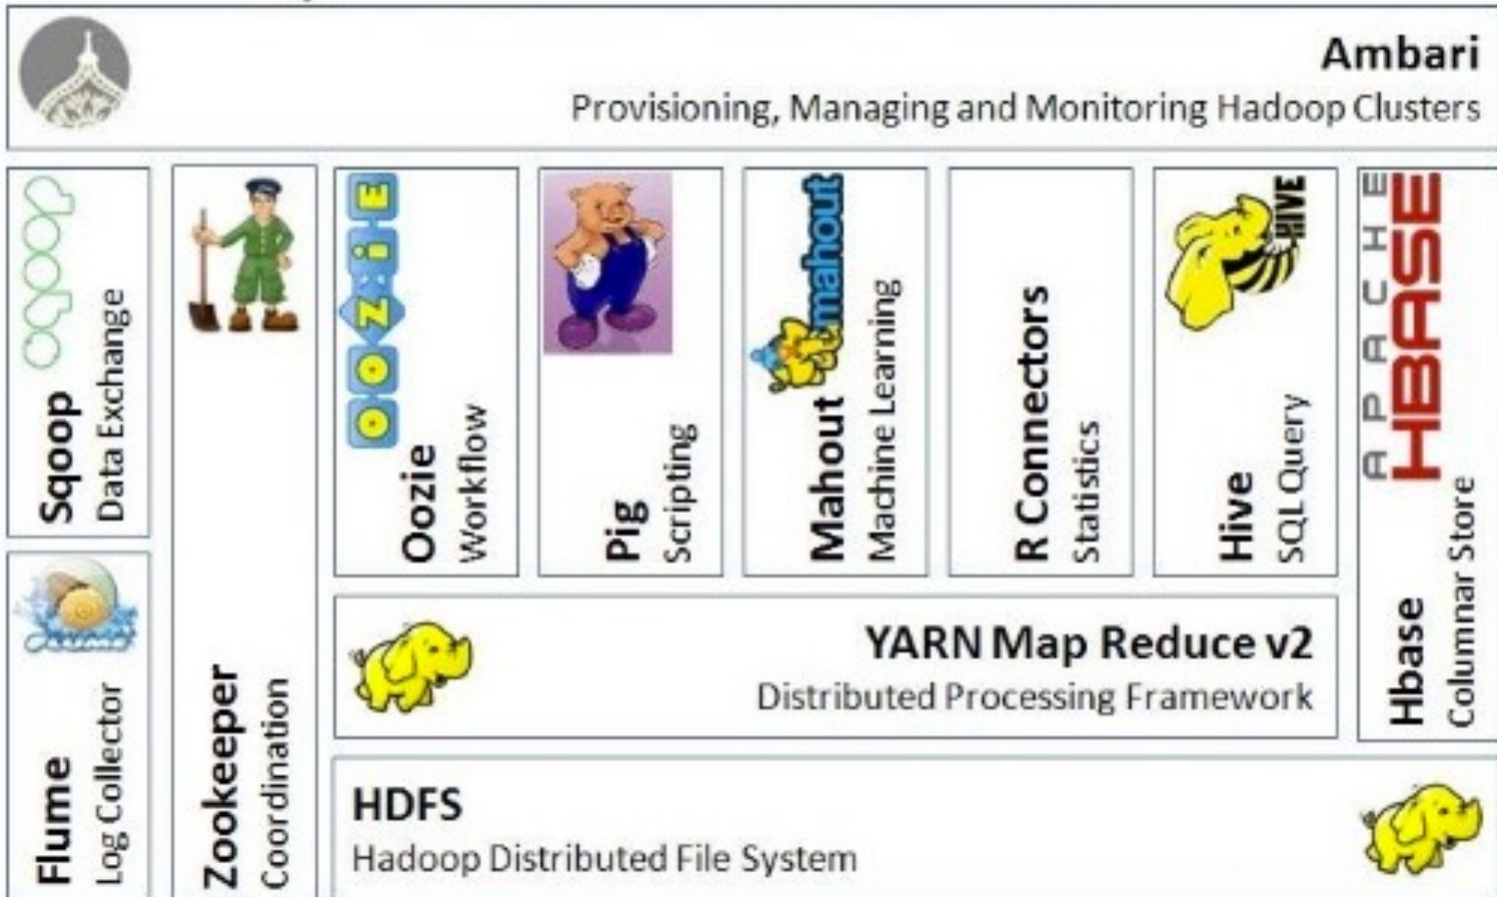
# Hadoop(Continued..)

- Hadoop is a platform that provides both distributed storage and computational capabilities.
- It was built mainly as a ***batch processing*** system.
- It consist of one cycle of loading data from HDFS applying Map job then running the reduce job and finally writing the output back into HDFS.
- In case of ***iterative*** computation the data needs to be loaded again!

# Hadoop EcoSystem



## Apache Hadoop Ecosystem





UNIVERSITAT DE  
BARCELONA

# HDFS

# Some Use case

- In 2010, **Facebook** claimed to have one of the largest HDFS cluster storing **21 Petabytes** of data.
- In 2012, **Facebook** declared that they have the largest single HDFS cluster with more than **100 PB** of data.
- And **Yahoo!** has more than **100,000 CPU** in over **40,000 servers** running Hadoop, with its biggest Hadoop cluster running **4,500 nodes**. All told, Yahoo! stores **455 petabytes** of data in HDFS.





# HDFS?

Hadoop Distributed file system or HDFS is a Java based distributed file system that allows you to store large data across multiple nodes in a Hadoop cluster.

It is based out of commodity hardware.


One of the main advantage is **Data Locality**

# Data Locality

Data locality talks about moving processing unit to data rather than the data to processing unit. In our traditional system, we used to bring the data to the application layer and then process it. But now, because of the architecture and huge volume of the data, bringing the data to the application layer will reduce the network performance to a noticeable extent. So, in HDFS, we bring the computation part to the data nodes where the data is residing. Hence, you are not moving the data, you are bringing the program or processing part to the data.

# HDFS Architecture

- **Apache HDFS** is a block-structured file system where each file is divided into blocks of a pre-determined size. These blocks are stored across a cluster of one or several machines. Apache Hadoop HDFS Architecture follows a *Master/Slave Architecture*, where a cluster comprises of a single NameNode (Master node) and all the other nodes are DataNodes (Slave nodes).

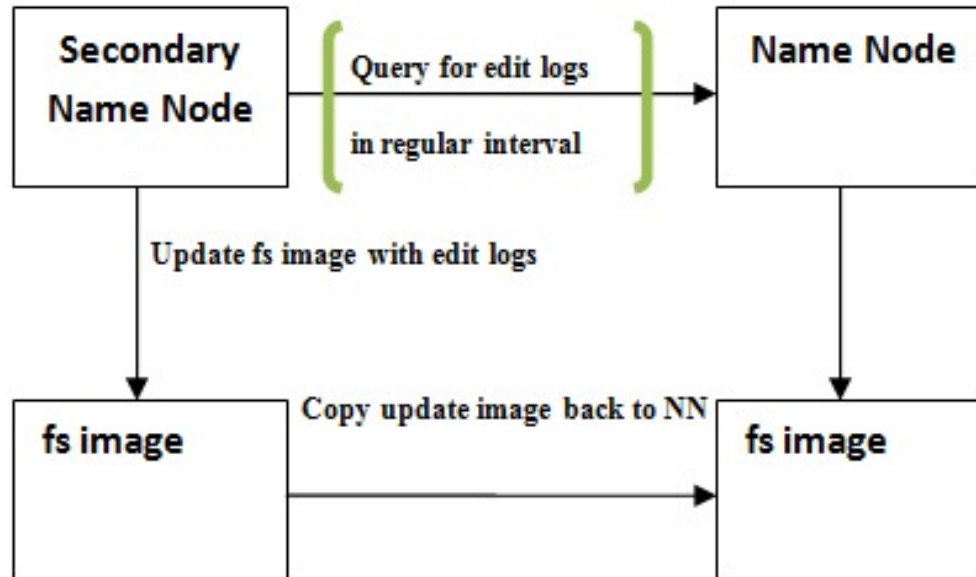


# HDFS Architecture

- **NameNode:** NameNode can be considered as a **master of the system**. It maintains the **file system tree** and the **metadata for all the files and directories present in the system**. Two files '**Namespace image**' and the '**edit log**' are used to store metadata information. Namenode has knowledge of all the datanodes containing data blocks for a given file, however, it does not store block locations persistently. This information is reconstructed every time from datanodes when the system starts.
- **DataNode:** DataNodes are **slaves** which reside on each **machine in a cluster** and **provide the actual storage**. It is responsible for serving, read and write requests for the clients.

# Secondary NameNode

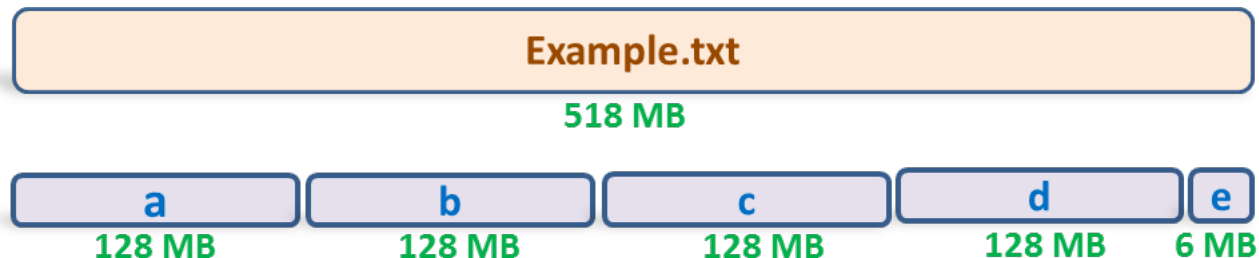
The Secondary NameNode works concurrently with the primary NameNode as a **helper daemon**. **It is not a backup of NameNode but just a helper.**




Secondary NameNode performs regular checkpoints in HDFS. Therefore, it is also called CheckpointNode.

# Blocks

Blocks are the nothing but the smallest continuous location on your hard drive where data is stored. In general, in any of the File System, you store the data as a collection of blocks. Similarly, HDFS stores each file as blocks which are scattered throughout the Apache Hadoop cluster. The default size of each block is 128 MB in Apache Hadoop 2.x (64 MB in Apache Hadoop 1.x) which you can configure as per your requirement.



# Replication management



Replication Factor 3



UNIVERSITAT DE  
BARCELONA

# Hadoop Processing





# Lets see how distributed computing work

Counting words in a document ( *You might ask why we would do such thing 😊 !! But it's a key step in Document search and matching*)

## One simple approach

## Use a Hash Table

# “Who is Ram

# Ram is who

# Ram is king

## King of Ajodhya”

[illegible]

## One simple approach

## Use a Hash Table

# “Who is Ram

# Ram is who

# Ram is king

## King of Ajodhya”

[illegible]

## King of Ajodhya”

| Key | Value |
|-----|-------|
| Who | 1     |
| Is  | 1     |
|     |       |
|     |       |
|     |       |
|     |       |
|     |       |

# One simple approach

Use a Hash Table

“Who is **Ram**

Ram is who

Ram is king

King of Ajodhya”

| Key | Value |
|-----|-------|
| Who | 1     |
| Is  | 1     |
| Ram | 1     |
|     |       |
|     |       |
|     |       |
|     |       |

# One simple approach

Use a Hash Table

“Who is Ram

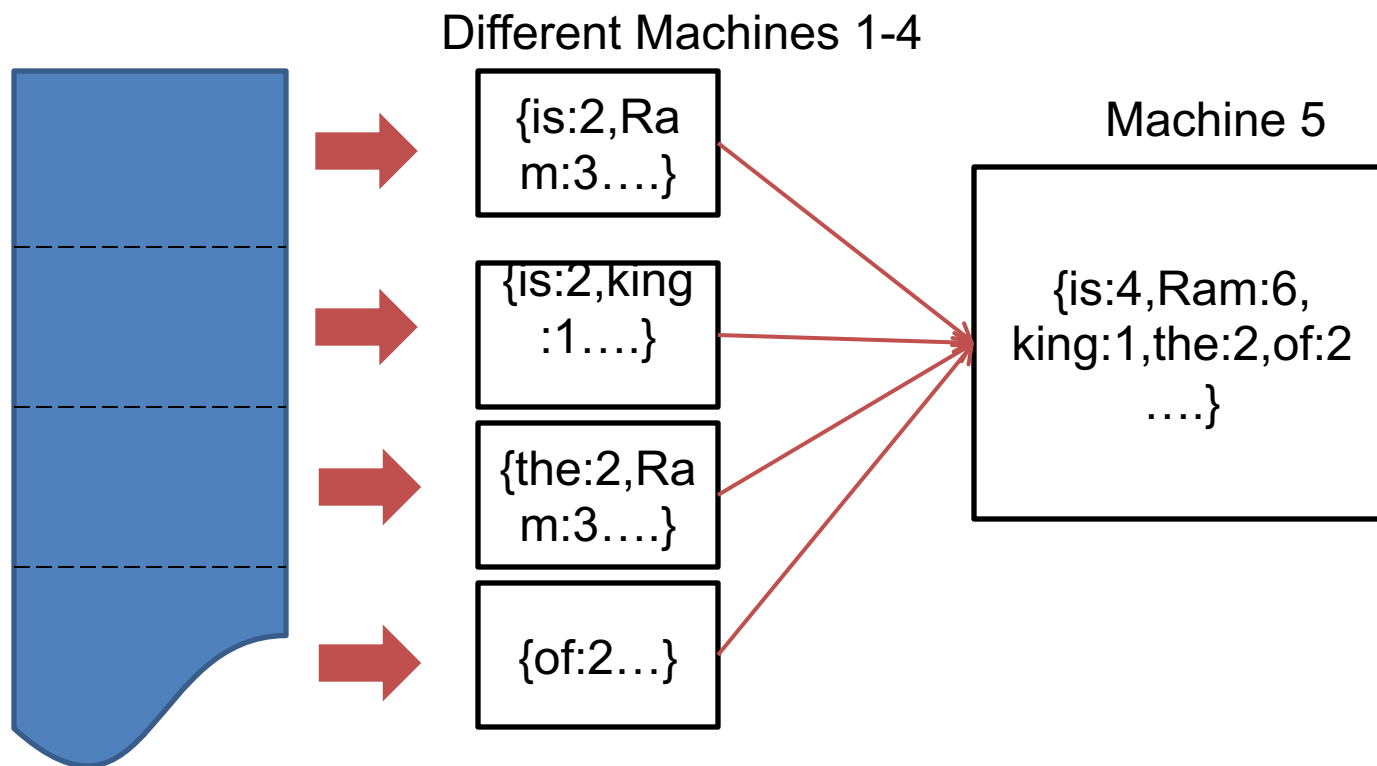
Ram is who

Ram is king

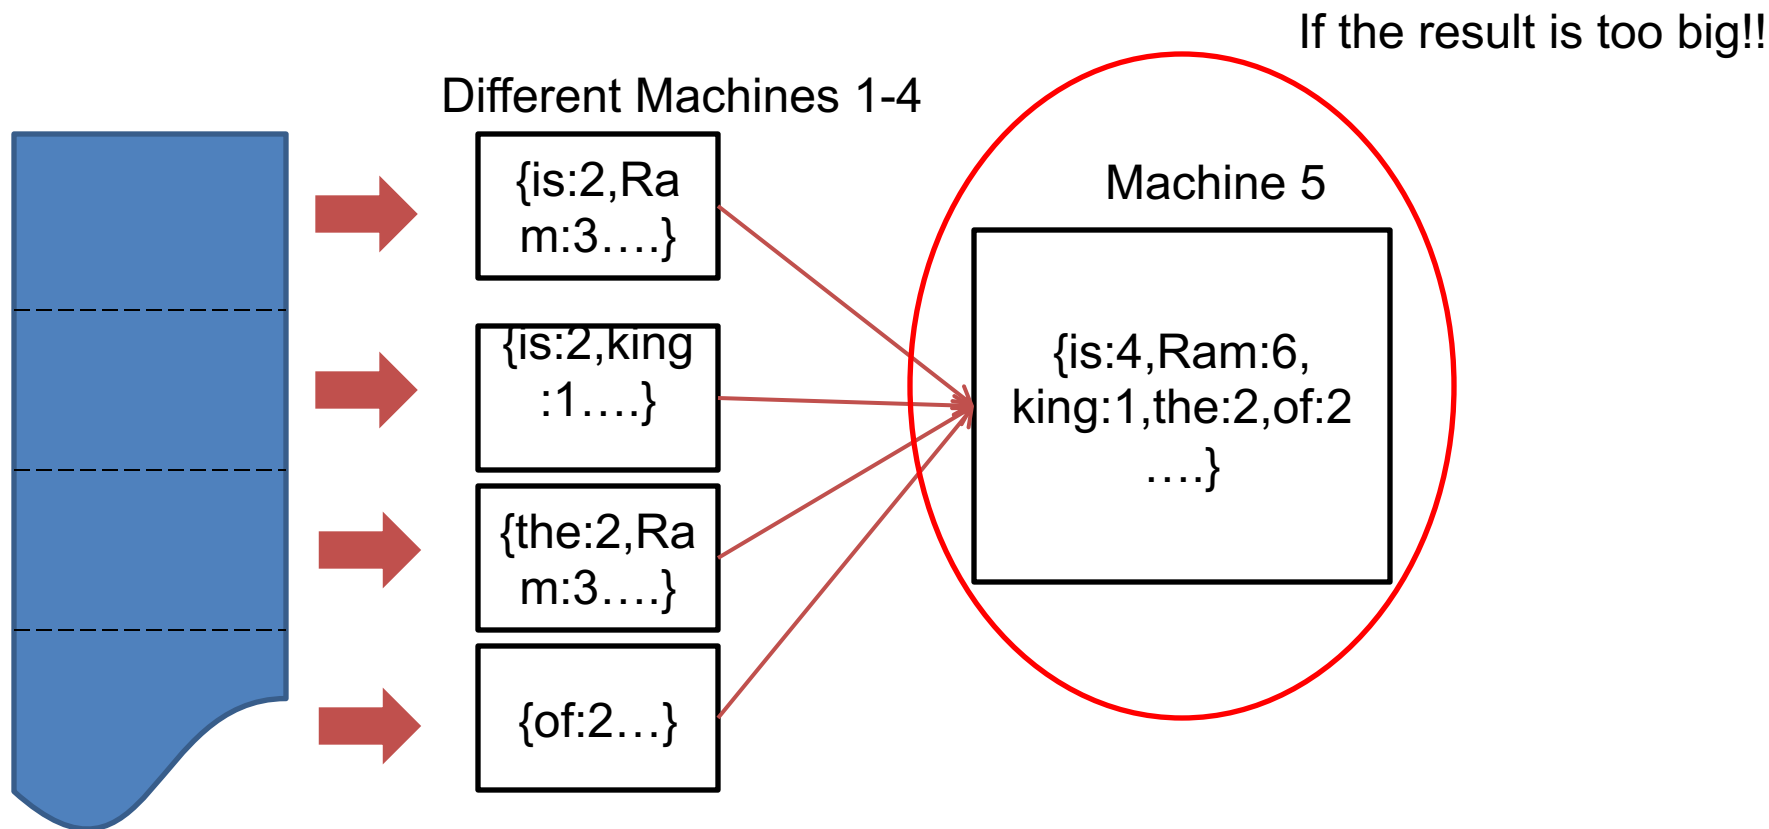
King of Ajodhya”

| Key | Value |
|-----|-------|
| Who | 1     |
| Is  | 1     |
| Ram | 2     |
|     |       |
|     |       |
|     |       |
|     |       |

# What if the document is very big?

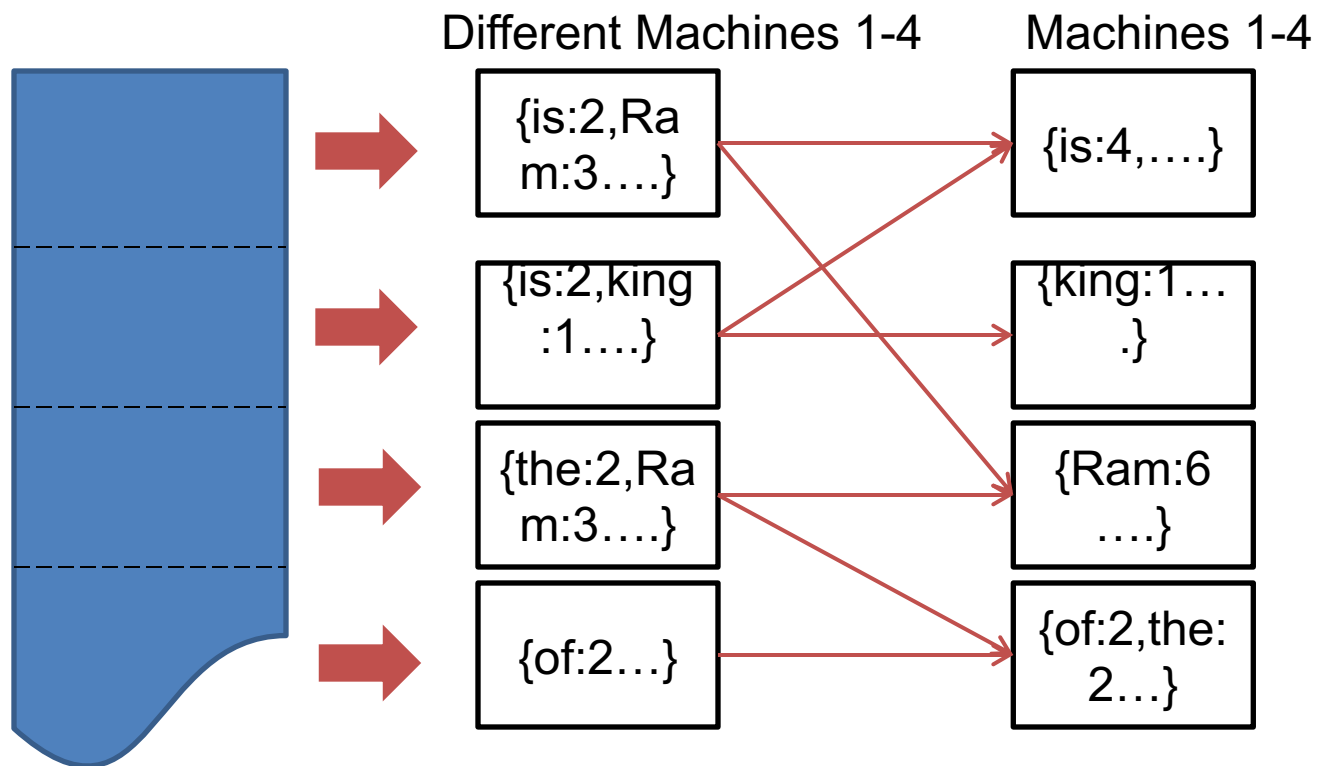


# What if the document is very big?

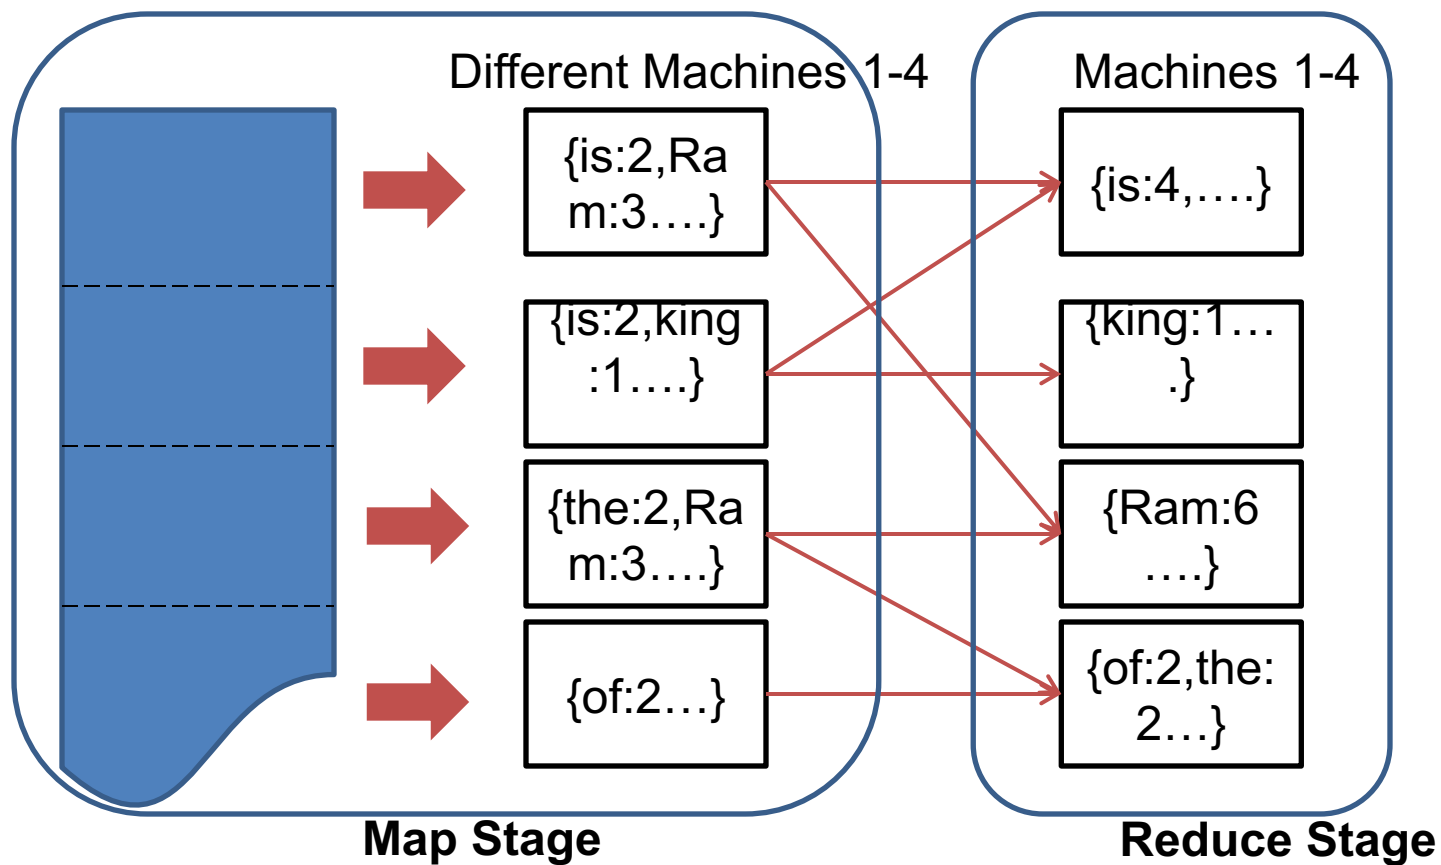




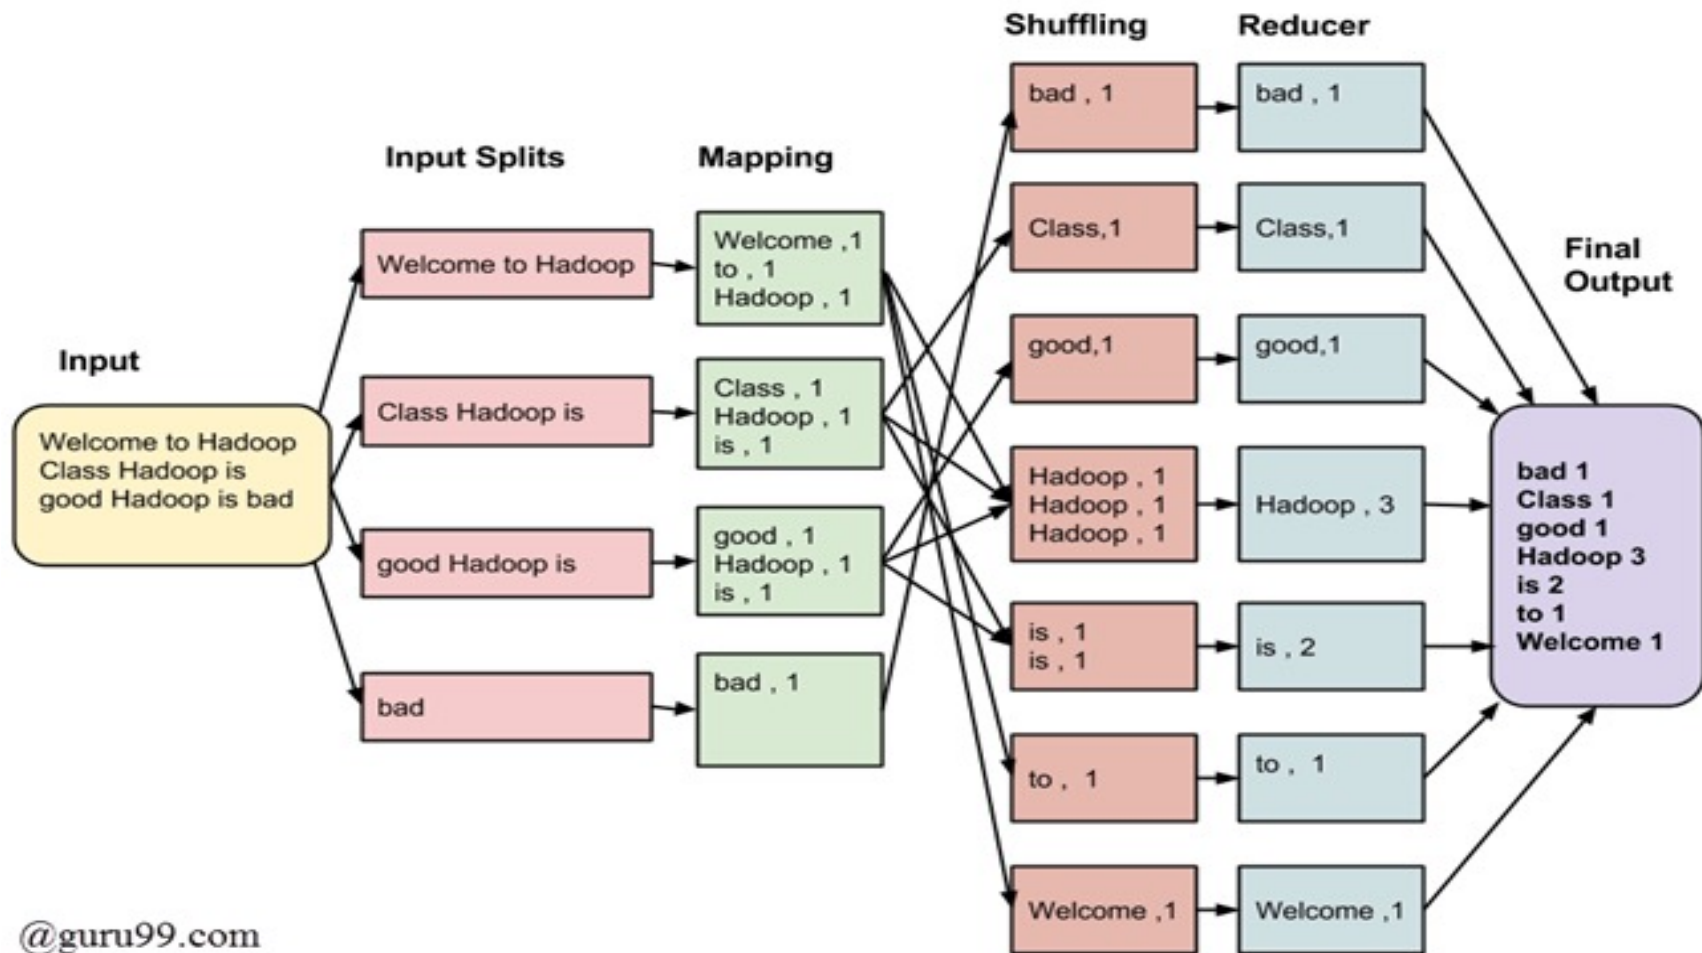
# Map Reduce



# Map Reduce



# Different Phases in MapReduce



@guru99.com

# MapReduce Phases

- **Input Splits:** An input to a MapReduce job is divided into fixed-size pieces called **input splits**. An input split is a chunk of the input that is consumed by a single map.
- **Mapping:** This is the very first phase in the execution of a map-reduce program. In this phase, data in each split is passed to a mapping function to produce output values. In our example, a job of the mapping phase is to count the number of occurrences of each word from input splits (more details about input-split are given below) and prepare a list in the form of <word, frequency>.

# MapReduce Phases

- **Shuffling:** This phase consumes the output of Mapping phase. Its task is to consolidate the relevant records from Mapping phase output. In our example, the same words are clubed together along with their respective frequency.
- **Reducing:** In this phase, output values from the Shuffling phase are aggregated. This phase combines values from Shuffling phase and returns a single output value. In short, this phase summarizes the complete dataset.

# Hadoop 2.0



**Hadoop v1.0**

**MapReduce**

Data Processing  
& Resource Management

**HDFS**

Distributed File Storage

2006-2012



**Hadoop v2.0**

**MapReduce**

Other Data  
Processing  
Frameworks

**YARN**

Resource Management

**HDFS**

Distributed File Storage

2013

# Read at home

- HDFS (Important)
- Watch <https://www.youtube.com/watch?v=ziqx2hJY8Hg>
- MapReduce
- Watch <https://www.youtube.com/watch?v=vbi95iqsnnM>